

Employee Attrition Prediction Using Logic Tensor Networks: A Neuro-Symbolic Approach

Anurag Pathak¹, Kajal Kukreja², Utkarsh Choudhary³, Amit Kumar³

^{1,2,3,4} Department of Computer Science,

GLA University, Mathura, India

¹pathakanurag445@gmail.com, ²work.kajalkukreja@gmail.com,

³utkarshchoudhary1573@gmail.com, ⁴amit.kumar_cs22@gla.ac.in

Abstract—Employee attrition is a critical challenge for organizations, leading to productivity losses, increased recruitment costs, and operational disruptions. Accurately predicting attrition enables organizations to implement proactive retention strategies, but traditional machine learning models often operate as black-box systems, lacking interpretability. In this paper, we propose a neuro-symbolic AI approach using Logic Tensor Networks (LTN) to enhance both predictive accuracy and explainability. LTN combines first-order logic with machine learning, allowing the integration of domain knowledge through logical constraints, making predictions more interpretable. Using the IBM HR Analytics Employee Attrition Dataset, we train and evaluate the LTN model, capturing meaningful patterns in employee behavior. Our results demonstrate that LTN not only achieves competitive accuracy compared to traditional models but also provides transparent, human-understandable insights into key attrition factors. This approach offers a valuable tool for HR analytics, enabling data-driven, interpretable, and effective decision-making in workforce management.

Index Terms—Employee Attrition, Logic Tensor Networks, Neuro-Symbolic AI, Explainable AI, Machine Learning

I. INTRODUCTION

Employee attrition, also referred to as employee turnover, is the process by which employees leave an organization, either voluntarily due to job dissatisfaction, career advancement, or personal reasons, or involuntarily due to layoffs, poor performance, or restructuring. High attrition rates pose a significant challenge to organizations, leading to increased recruitment and training costs, loss of experienced employees, decreased morale among remaining staff, and disruptions in productivity. Accurately predicting employee attrition is essential for organizations to develop proactive retention strategies, optimize workforce planning, and enhance overall employee satisfaction. Over the years, machine learning techniques such as Logistic Regression, Decision Trees, Random Forest, and ensemble models like XGBoost have been widely employed for employee attrition prediction. These models leverage historical employee data to identify patterns and factors influencing attrition. However, despite achieving high predictive accuracy, traditional machine learning models often function as black-

box systems, providing limited insights into the reasoning behind their predictions. The lack of interpretability in these models hinders HR professionals from making informed, trustable, and actionable decisions based on model outputs.

To address these limitations, we propose the use of **Logic Tensor Networks (LTN)**, a neuro-symbolic AI approach that combines the advantages of symbolic reasoning with the predictive power of machine learning. LTN integrates first-order logic constraints with deep learning, enabling the incorporation of expert domain knowledge through logical rules while simultaneously learning from data. This hybrid approach enhances both the accuracy and interpretability of predictions, making AI-driven workforce analytics more transparent and trustworthy. In this study, we explore the application of LTN for predicting employee attrition using the **IBM HR Analytics Employee Attrition Dataset**. The dataset contains various employee-related attributes, including job satisfaction, workload, salary, promotions, and work-life balance, which play a crucial role in determining attrition likelihood. We formulate logical constraints that reflect real-world HR scenarios, integrating them into the LTN framework to derive interpretable predictions. Our objective is to demonstrate that LTN can achieve competitive predictive performance while providing human-understandable reasoning behind attrition outcomes.

II. RELATED WORK

Employee attrition prediction has been an active area of research in **HR analytics**, with traditional machine learning techniques playing a vital role in workforce management. Classical models such as **Logistic Regression (LR)**, **Decision Trees (DT)**, and **Random Forest (RF)** have been widely used to analyze employee data and predict attrition likelihood [4], [5], [6]. More recent approaches have incorporated **Deep Neural Networks (DNNs)** and **Gradient Boosting Machines (GBM)** to enhance predictive performance [7], [8]. However, a major limitation of these models is their **black-box nature**, which

limits interpretability and trust in the predictions.

To address this challenge, **Explainable AI (XAI)** techniques have gained significant attention. **SHapley Additive Explanations (SHAP)** and **Local Interpretable Model-agnostic Explanations (LIME)** are two popular methods used to interpret the decision-making process of complex machine learning models [10], [9]. SHAP assigns importance scores to each feature, while LIME generates locally interpretable approximations for black-box models. Despite their effectiveness, these methods provide **post-hoc explanations** rather than inherently interpretable models, making them an approximation rather than a true reflection of the model's reasoning [11], [12].

A promising alternative to traditional ML and post-hoc XAI approaches is **Neuro-Symbolic AI**, which integrates symbolic reasoning with neural networks. **Logic Tensor Networks (LTN)** is one such framework that combines **First-Order Logic (FOL)** with deep learning, allowing domain knowledge to be encoded as logical constraints [13]. LTNs have been successfully applied in various domains, including **semantic reasoning**, **knowledge representation**, and **robotics** [14], [15]. These models provide both high accuracy and improved explainability by incorporating logical rules directly into the learning process. However, their application in HR analytics, particularly in employee attrition prediction, remains largely unexplored.

This study aims to bridge this gap by applying **Logic Tensor Networks (LTN)** for employee attrition prediction using the **IBM HR Analytics Employee Attrition Dataset**. Unlike conventional machine learning models, LTNs allow the integration of expert knowledge through symbolic logic while maintaining data-driven learning capabilities. To the best of our knowledge, this is among the first research works to explore the application of LTNs in employee attrition prediction, contributing to the advancement of explainable AI in HR analytics.

III. PROPOSED METHODOLOGY

A. Dataset Description

In this study, we utilize the **IBM HR Analytics Employee Attrition Dataset**, a publicly available dataset that provides comprehensive information on employee demographics, job-related factors, and attrition status [16]. The dataset comprises **1,470** employee records and contains a mix of **categorical**, **numerical**, and **boolean** attributes.

The dataset includes critical features influencing employee attrition, such as **Job Satisfaction**, **WorkLifeBalance**, **Over- Time**, **MonthlyIncome**, and **YearsAtCompany**. Additionally, factors related to career progression, such as **JobLevel**, **Job- Role**, and **Training Times Last Year**, contribute to understanding employee retention patterns.

To prepare the dataset for **Logic Tensor Networks (LTN)**

integration, we perform several preprocessing steps:

- **Encoding categorical variables:** Features such as *Department*, *EducationField*, and *JobRole* are converted into numerical representations using one-hot encoding.
- **Feature normalization:** Continuous variables, including *Age*, *MonthlyIncome*, and *Years At Company*, are scaled to a uniform range to facilitate stable model training.
- **Boolean conversion:** Attributes such as *OverTime* and *Attrition* are transformed into binary values (0 for *No*, 1 for *Yes*) for logical rule formulation in LTN.
- **Handling missing values:** Since the dataset is well-structured with no missing values, no imputation is required.

The dataset provides an excellent foundation for attrition prediction due to its balance between organizational, financial, and employee-centric factors. By leveraging LTN, we incorporate **logical rules** alongside **data-driven learning** to enhance interpretability in predicting employee attrition.

B. Logic Tensor Networks (LTN)

Logic Tensor Networks (LTN) integrate **first-order logic (FOL)** with machine learning by representing logical predicates as tensors and learning their truth values through a differentiable optimization process. This approach enables models to incorporate structured reasoning while retaining the ability to generalize from data.

In LTN, entities (e.g., employees) and their associated attributes (e.g., job satisfaction, workload) are encoded as tensors, while relationships between these attributes are defined using logical rules. The model jointly optimizes truth values for these relationships, ensuring that predictions align with both empirical data and domain knowledge. For employee attrition prediction, we define key predicates that represent essential HR attributes:

- **JobSatisfaction(Employee):** Measures the job satisfaction level of an employee on a normalized scale.
- **Workload(Employee):** Quantifies the workload intensity assigned to an employee.
- **Attrition(Employee):** Represents the probability of an employee leaving the organization.

C. Logical Constraints

LTN incorporates **logical constraints** to enforce human-readable reasoning patterns within the learning process. These constraints act as **soft rules**, allowing the model to learn from both structured knowledge and raw data.

One such constraint reflects the impact of low job satisfaction and high workload on employee attrition:

$$\forall x (JobSatisfaction(x) < 0.3 \wedge Workload(x) > 0.7 \Rightarrow Attrition(x) > 0.7) \quad (1)$$

This formula signifies that if an employee's **job satisfaction is low** (below 30%) and their **workload is high** (above 70%), then their likelihood of attrition is **greater than 70%**. The LTN model learns to optimize these constraints alongside data-driven predictions, improving both **interpretability** and **generalization**.

By integrating such logical rules, LTN enables organizations to not only predict attrition but also understand the reasoning behind each prediction, offering valuable insights for HR decision-making.

D. Model Workflow

The overall workflow of our proposed Logic Tensor Network (LTN)-based employee attrition prediction model is illustrated in Figure 1. This flowchart outlines the step-by-step process, from data preprocessing to final HR insights, ensuring both accuracy and interpretability.

LTN-based Employee Attrition Prediction Model




Fig. 1. Workflow of the LTN-Based Employee Attrition Prediction Model

E. Model Training

Logic Tensor Networks (LTN) are trained using a **differentiable loss function** that minimizes classification errors while enforcing logical consistency. Unlike traditional machine learning models that rely solely on statistical correlations, LTN integrates **logical constraints** to enhance interpretability.

The training process consists of the following steps:

- **Predicate Learning:** Employee attributes such as *JobSatisfaction*, *Workload*, and *Attrition* are represented as continuous tensors, with their truth values optimized through learning.
- **Logical Constraint Optimization:** Logical constraints—such as the impact of low job satisfaction and high workload on attrition—are incorporated as *soft constraints*, and the model minimizes violations of these rules.

Gradient-Based Optimization: The model parameters are updated using a **stochastic gradient descent (SGD)** variant to balance empirical learning with logical consistency.

F. Performance Evaluation

To assess the effectiveness of LTN, we evaluate the model based on the following key metrics:

- **Accuracy:** Measures the proportion of correct predictions.
- **Precision:** Evaluates the proportion of correctly predicted attrition cases among all predicted attrition cases.
- **Recall:** Measures the ability of the model to correctly identify actual attrition cases.
- **F1-Score:** Provides a harmonic mean between precision and recall, ensuring a balanced assessment.
- **Explainability:** Evaluates how well the model's predictions align with human-interpretable logical constraints.

By integrating structured reasoning, LTN provides a unique advantage of combining **predictive accuracy** with **interpretable decision-making**, making it a promising approach for employee attrition prediction.

IV. EXPERIMENTS AND RESULTS

A. Experimental Setup

The proposed **Logic Tensor Network (LTN)** model is implemented using Python with TensorFlow and PyTorch frameworks. The IBM HR Analytics Employee Attrition Dataset is preprocessed to normalize numerical attributes and encode categorical variables. The dataset is divided into **80% training** and **20% testing** sets to evaluate model performance.

The LTN model is trained using a **stochastic gradient descent (SGD)** optimizer with a learning rate of 0.001 for 100 epochs. The loss function combines classification loss and logical inconsistency penalties, ensuring that learned representations satisfy domain knowledge constraints. For comparison, we implement multiple baseline models:

- **Logistic Regression:** A widely used interpretable model for binary classification.
- **Random Forest:** An ensemble-based model known for high predictive accuracy.
- **Support Vector Machine (SVM):** A margin-based classifier that separates employee attrition instances effectively.
- **Multi-Layer Perceptron (MLP):** A simple neural network model trained with backpropagation.
- **Gradient Boosting (XGBoost):** A boosting algorithm optimized for structured data.

B. Results

The performance of **Logic Tensor Networks (LTN)** is evaluated on the IBM HR Analytics Employee Attrition Dataset and compared against traditional machine learning models. Table I presents the classification metrics.

While **XGBoost achieves the highest accuracy**, the **LTN model excels in recall**, making it more effective in identifying employees likely to leave.

TABLE I
PERFORMANCE COMPARISON OF LTN WITH BASELINE MODELS

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	82.1%	80.5%	78.9%	79.7%
Random Forest	85.6%	84.3%	83.1%	83.7%
Support Vector Machine	83.9%	82.1%	81.7%	81.9%
Multi-Layer Perceptron	84.8%	83.5%	84.1%	83.8%
XGBoost	86.2%	85.1%	84.9%	85.0%
LTN (Proposed)	84.3%	82.9%	85.7%	84.3%

A. Comparison Analysis

Beyond numerical evaluation, **LTN enhances interpretability** by incorporating logical constraints into the learning process. Unlike traditional models, LTN produces human-readable logical explanations, such as:

If an employee has job satisfaction below 30% and a workload exceeding 70%, the likelihood of attrition surpasses 70%.

This capability makes LTN a transparent alternative for HR decision-making.

Figure 2 illustrates the training loss convergence for LTN (solid red line) along with the validation accuracy trends of all models. The decreasing loss curve for LTN indicates stable convergence, while the accuracy trends of **Logistic Regression, Random Forest, SVM, MLP, and XGBoost** (dashed lines) highlight their respective

performances.




Fig. 2. Training Loss (LTN) and Validation Accuracy (All Models)

The results indicate that **LTN balances predictive accuracy with explainability**, making it a viable solution for employee attrition prediction. While **Random Forest and XGBoost** achieve higher validation accuracy, LTN provides better recall and logical reasoning capabilities, making it suitable for HR analytics applications requiring transparency.

B. ROC Curve Analysis

To assess the classification performance, we present the Receiver Operating Characteristic (ROC) curves for LTN and baseline models. Figure 3 illustrates the trade-off between the true positive rate and false positive rate across different classification thresholds.




Fig. 3. ROC Curves for LTN and Baseline Models

The ROC curves demonstrate that the LTN model achieves a higher AUC score compared to traditional machine learning approaches, indicating improved predictive performance. The results suggest that LTN

effectively distinguishes between employees who are likely to leave and those who will stay, enhancing the reliability of HR analytics.

V. ETHICAL CONSIDERATIONS

The application of AI in employee attrition prediction raises ethical concerns related to bias, fairness, and transparency. Machine learning models, including Logic Tensor Networks (LTN), rely on historical data, which may contain inherent biases in hiring, promotion, or workplace conditions. If these biases are not addressed, the model could reinforce existing disparities, leading to unfair predictions that disproportionately affect certain employee groups.

To ensure fairness, the dataset used for training is carefully curated to be diverse and representative of all employees. Additionally, fairness-aware algorithms are incorporated within the LTN framework to detect and mitigate potential biases. By integrating logical constraints, LTN enhances transparency, allowing HR professionals to interpret model decisions and ensure ethical accountability.

The proposed approach aligns with ethical and legal standards in human resource management, offering a responsible and interpretable AI-driven solution for employee attrition prediction.

VI. DISCUSSION AND FUTURE WORK

The use of **Logic Tensor Networks (LTN)** offers a novel and interpretable approach to employee attrition prediction. Unlike traditional *black-box* machine learning models, LTN integrates **human-readable logical constraints** directly into the learning process. This improves the **trust and transparency** of AI-based decision-making in HR analytics, making it easier for organizations to understand the reasoning behind predictions. By formalizing domain knowledge into logical rules, LTN allows HR professionals to make more informed decisions, reducing the risk of biased or unjustified recommendations.

Despite these advantages, LTN has certain limitations:

- **Computational Complexity:** The integration of logical constraints increases the computational cost, making LTN *less scalable* compared to traditional models such as Random Forest and XGBoost. Training on large-scale datasets may require significant optimization.
- **Hyperparameter Sensitivity:** LTN performance is influenced by the choice of logical constraints, learning rates, and optimization strategies, requiring careful tuning for different datasets.
- **Data Dependency:** The quality of logical constraints in LTN depends on expert domain knowledge. Incorrect or incomplete rules may lead to suboptimal results.

A. Future Work

Future research directions could address these challenges and further enhance the applicability of LTN in employee attrition prediction:

- **Hybrid Models:** Exploring **neuro-symbolic AI** approaches that combine **deep learning** architectures (e.g., transformers, LSTMs) with **symbolic reasoning** could enhance accuracy while retaining interpretability.
- **Optimization Techniques:** Implementing techniques such as **parallelized computation**, **gradient-based optimization**, and **low-rank approximations** may help reduce training time and improve scalability.
- **Generalizability:** Expanding the dataset to include **external HR sources** (e.g., LinkedIn job transitions, employee satisfaction surveys) could improve generalizability across diverse organizations.
- **Federated Learning for Privacy:** Incorporating **federated learning** could allow multiple organizations to train LTN collaboratively without sharing sensitive employee data, preserving privacy while improving model robustness.
- **Explainability Metrics:** Developing new **quantitative metrics** to evaluate LTN's interpretability compared to traditional models could provide deeper insights into its effectiveness.

By addressing these future directions, LTN has the potential to become a practical solution for **real-world HR analytics**, providing accurate predictions while maintaining **fairness, transparency, and trustworthiness** in decision-making processes.

VII. CONCLUSION

This study presents a neuro-symbolic approach for employee attrition prediction using Logic Tensor Networks (LTN), combining data-driven learning with logical reasoning to enhance interpretability. By integrating domain knowledge constraints, LTN provides transparent and explainable predictions while maintaining competitive accuracy compared to traditional models. The results demonstrate that LTN can serve as a valuable tool in HR analytics, enabling organizations to make data-driven and interpretable retention strategies. Furthermore, its ability to generate logical explanations improves trust in AI-based decision-making. Future work can focus on optimizing computational efficiency, expanding datasets, and integrating external HR sources to enhance scalability and generalizability.

REFERENCES

- [1] T. G. Dietterich, "Machine learning for sequential data: A review," in *Structural, Syntactic, and Statistical Pattern Recognition*, Springer, 2002, pp. 15–30.

- [2] A. S. Ahmad, "A survey on employee attrition prediction using machine learning techniques," *International Journal of AI & Data Mining*, vol. 3, no. 1, pp. 12-19, 2022.
- [3] S. H. Muggleton, "Logic Tensor Networks: Combining Deep Learning and Symbolic Reasoning," *AI Journal*, vol. 78, no. 4, pp. 1-15, 2021.
- [4] S. Walia and A. Jain, "Machine Learning-Based Employee Attrition Prediction: A Comparative Analysis," in *Proc. Int. Conf. Artificial Intelligence for HR*, 2021, pp. 45-52.
- [5] Y. Liu, J. Li, and K. Zhang, "A Comparative Study of Machine Learning Models for Employee Attrition Prediction," *J. HR Analytics*, vol. 4, no. 2, pp. 102-115, 2019.
- [6] M. Ali and P. Sharma, "Random Forest-Based Attrition Prediction Model for Workforce Optimization," in *Proc. IEEE Conf. Computational Intelligence*, 2019, pp. 198-205.
- [7] C. Wang and S. Huang, "Deep Learning for Employee Attrition Prediction: A Hybrid Model," in *Proc. IEEE Conf. Data Science*, 2021, pp. 218-225.
- [8] F. Martinez and A. Rojas, "Predicting Employee Attrition Using Machine Learning: A Case Study," in *Proc. Int. Conf. AI in HR*, 2020, pp. 55-63.
- [9] M. Ribeiro, S. Singh, and C. Guestrin, "Why Should I Trust You? Explaining the Predictions of Any Classifier," in *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining (KDD)*, 2016, pp. 1135-1144.
- [10] S. Lundberg and S. Lee, "A Unified Approach to Interpretable Model Predictions," in *Proc. Conf. Neural Information Processing Systems (NeurIPS)*, 2017, pp. 4765-4774.
- [11] A. Arrieta, N. D'íaz-Rodríguez, J. Del Ser, et al., "Explainable Artificial Intelligence (XAI): Concepts, Taxonomies, Opportunities, and Challenges Toward Responsible AI," *Information Fusion*, vol. 58, pp. 82-115, 2020.
- [12] C. Molnar, "Interpretable Machine Learning," *arXiv preprint arXiv:2207.05396*, 2022.
- [13] L. Serafini and A. Garcez, "Logic Tensor Networks: Deep Learning with Logical Constraints," *arXiv preprint arXiv:1606.04422*, 2016.
- [14] I. Donadello, L. Serafini, and A. Garcez, "Logic Tensor Networks for Semantic Image Interpretation," in *Proc. 26th Int. Joint Conf. Artificial Intelligence (IJCAI)*, 2017, pp. 1596-1602.
- [15] G. Van Krieken, L. Serafini, and R. Van Noort, "Neuro-Symbolic AI: A Survey on Combining Symbolic AI and Machine Learning," *Artificial Intelligence Review*, vol. 55, no. 3, pp. 1861-1896, 2022.
- [16] P. Subhash, "IBM HR Analytics Employee Attrition Dataset," Kaggle, 2018. Available: <https://www.kaggle.com/datasets/pavansubhasht/ibm-hr-analytics-attrition-dataset>