

Hybrid Deep Learning Model for Social Media Sentiment Analysis

1st C.Hrishikesava Reddy
department of CSE

Rajeev Gandhi Memorial College of
Engineering and Technology
Nandyal, Andhra Pradesh-
518501
hrushi.c@rgmcet.edu.in

2nd M. Durga Prasad
department of CSE

Rajeev Gandhi Memorial College of
Engineering and Technology
Nandyal, Andhra Pradesh-
518501
durgaprasad312003@gmail.co
m

3rd S. Dhanush Kumar Reddy
department of CSE

Rajeev Gandhi Memorial College of
Engineering and Technology
Nandyal, Andhra Pradesh-
518501
s.dhanush1375@gmail.com

4th k. Giri Babu
department of CSE

Rajeev Gandhi Memorial College of
Engineering and Technology
Nandyal, Andhra Pradesh-
518501
giribabu9059@gmail.com

5th S. Jayanth
department of CSE

Rajeev Gandhi Memorial College of
Engineering and Technology
Nandyal, Andhra Pradesh-
518501
Sunkarijayanth17@gmail.com

Abstract: The informal language, sarcasm and context-dependent expressions, which are informal, are a big challenge to the sentiment analysis of the social media content because the traditional models cannot deal with them. The paper describes a context-specific sentiment analysis system, that is, three forms of hybrid deep learning: BiLSTM, RoBERTa-BiLSTM, and DistilBERT +GRU to categorize aspects to sentiment on conventional Imdb and twitter data. The models are appropriately designed to mimic local syntactic regularities, and global syntactic dependencies of the noise and user generated text. To solve the black-box of deep-learning, we combine Explainable AI through the employment of LIME which gives visual explanations of any prediction. The framework is implemented as an interactive web-based application to carry out sentiment analysis in real-time through transparency reporting. RoBERTa-BiLSTM also becomes more precise on binary IMDB reviews with 90 percent accuracy and three-class Twitter data with 80 percent accuracy than single transformer models through experimental analysis. Competitive (89% IMDB, 79% Twitter) but with much less computation is DistilBERT +GRU. It represents a useful source in the procedure of social media monitoring, customer feedback analysis, and the monitoring of the opinion of the population since it is based on the state-of-the-art NLP techniques, besides the fact that it can be easily understood and user-friendly.

Keywords: *Sentiment Analysis, Aspect-Based Sentiment Analysis, Context-Aware Models, Explainable AI, LIME, Hybrid Deep Learning, Social Media Analytics, IMDB, Twitter, NLP.*

I. INTRODUCTION

Sentiment analysis has been of paramount importance in the extraction of information on user-generated contents on websites such as IMDB and Twitter where sentiments are passed loosely in an indiscreet and casual manner, with sarcasm and contingency. This information is a liability to

the old sentiment analysis systems especially in short noisy texts where sentiment is not explicitly known. The conversation that occurs in social media is very casual and this implies that the use of slang expressions, shortenings, and ambiguous phrases occurs which makes the task of categorizing the sentiments even more difficult in the traditional sense.

The new methods of deep learning have delivered spectacular results in the understanding of contextual language. Nevertheless, these models may be black box models very frequently or are projections that do not involve the disclosure of what is going on behind the scenes in their decision processes. The lack of clarity is a massive disadvantage to the trust of the user, particularly in such a sensitive application as the power to track the opinion of the people and customer response feedback.

The current sentiment analysis systems cannot work with contextual expression of such expressions occurring on user generated information at sites such as IMDB and Twitter. They are mostly short, non-formal and sarcastic or colloquial and this aspect makes bag-of-words and simple neural models pointless. Social media communication has a great deal of noise that encompasses the cues of sentiment, and high-level knowledge is dependency covered by the high level of contextuality as compared to the surface level that is largely covered by the words. Moreover, the majority of deep learning approaches are black-boxes that do not provide much data as to how they are making their decision. This creates poor transparency,

which minimizes the user confidence on sensitive application.

To help mitigate these defects, this paper will propose a context sentimental analysis model into three hybrid deep learning models that comprise BiLSTM, RoBERTa-BiLSTM and DistilBERT + GRU. These models have been trained on both IMDb and Twitter data to obtain both local syntactic and global semantic dependencies. Some of the key functions include ensuring models are made more transparent with the assistance of Explainable AI through LIME to be capable of possessing an explanation of individual predictions in visual form. It is a real-time sentiment analysis, and is presented as a Web application, which can be easily accessed.

The key findings of the specified paper are the following: (1) the comparative analysis of three hybrid architectures on formal and informal texts, (2) the added value of explainability through the implementation of LIME to be in a position to make transparent choices, (3) the analysis of sentiments in real-time, and (4) the proof of the fact that RoBERTa-BiLSTM has a higher accuracy (90% on IMDb, 80% on Twitter) and can be explained. The accuracy, transparency, and the useability of this framework makes it an effective solution to the marketing and customer feedback analysis trends and social media, thus, closing the gap that exists between advanced NLP algorithms and the readily available and usable sentiment analysis software.

II. RELATED WORKS

Sentiment analysis is an important instrument of natural language processing that involves delicate classification to derive information from user-generated content on websites such as IMDb and Twitter. The main concerns include dealing with informal language, sarcasm, ambiguity, and the black-box nature of deep learning models, as well as ensuring robustness across different datasets. Recent IEEE publications (2023–2026) reflect progress from traditional machine learning into hybrid models, transformers, multimodal fusion, and explainable AI methods for applications in e-commerce, finance, and mental health monitoring.

[1] presented a survey on explainable sentiment analysis, integrating XAI with sentiment models to address black-box issues, emphasizing local and global explanations for user trust. [2] introduced a multimodal sentiment analysis framework with a new taxonomy, exploring alignment and large language models like GPT. [3] reviewed e-commerce sentiment analysis, achieving up to 93% accuracy with SVM and highlighting aspect-based analysis and sarcasm detection. [4] proposed a systematic review of low-resource sentiment analysis, identifying transfer learning and BERT as dominant approaches with accuracies above 80%. [5] reviewed aspect-based

sentiment analysis, achieving 90% precision with rule-based methods and identifying gaps in implicit aspects. [6] surveyed challenges like cross-domain and multimodal data, emphasizing F1-score for reliable classification. [7] developed ClimateBERT for climate change tweets, achieving 90.22% accuracy. [8] reviewed LLMs in financial sentiment analysis, identifying interpretability gaps. [9] reviewed Twitter sentiment analysis, emphasizing hybrid models for crisis detection. [10] surveyed multimodal fusion methods, showing improved accuracy over unimodal approaches. [11] developed BERT-Fuse for mental health, achieving 97.05% accuracy. [12] proposed a recommendation system using sentiment analysis and matrix factorization, improving F1-scores by up to 19.2%.

Even with a considerable progress, there are still some gaps in research that have not been filled. To begin with, although there have been studies of hybrid models that involve transformers with recurrent models, few studies have made a systematic comparison of the various hybrid modes across formal and informal textual genres. Second, literature on the subject does not have unified frameworks that can provide high classification accuracy and explainable visual interfaces. Third, the majority of explainable AI usage in sentiment analysis is hypothetical, and not many are implemented as working web-based systems to allow real-time user engagement. To fill these gaps, this paper suggests a comparative analysis of three hybrid architectures, incorporation of LIME-based explainability, and implementation of the system as an interactive web application. Its novelty consists in the fact that RoBERTa-BiLSTM can reach higher accuracy (90% on IMDb, 80% on Twitter) and at the same time stay transparent, thus closing the gap between sophisticated tools in NLP and the tools that are useful and easily accessible to users.

Table 1: Comparison with Existing Methods

<u>Existing Method</u>	<u>Dataset</u>	<u>Reported Accuracy</u>	<u>Our Accuracy</u>	<u>Improvement</u>
Standard LSTM [9]	Twitter	75-78%	80%	+2-5%
Standard BERT [4]	Twitter	76-78%	80%	+2-4%
TF-IDF + SVM [3]	Twitter	72-75%	80%	+5-8%
Standard RoBERTa [7]	IMDb	87-88%	90%	+2-3%
Standard BERT [11]	IMDb	85-87%	90%	+3-5%
Baseline BiLSTM	IMDb	84%	90%	+6%
Baseline BiLSTM	Twitter	77%	80%	+3%

Table 1 presents a comparative analysis of our proposed RoBERTa-BiLSTM model against existing sentiment analysis methods reported in recent literature. The results demonstrate that our hybrid approach consistently outperforms standalone architectures across both datasets,

with accuracy improvements ranging from 2% to 8%. The most significant gains are observed on Twitter data, where our model achieves 80% accuracy compared to 72-75% with traditional machine learning approaches like TF-IDF + SVM. These improvements validate the effectiveness of combining transformer-based contextual understanding with sequential modeling for handling noisy, informal social media text.

III. PROPOSED SYSTEM




Figure 1: Flow diagram of the proposed system

The logic behind the proposed system Figure 1 is to process a user-generated text in social media, which is inputted into the context-sensitive sentiment analysis system, based on deep learning hybrid structures. The system has applied three other models such as BiLSTM, RoBERTa-BiLSTM, and DistilBERT +GRU models to find out syntactic and semantic properties of informatives that are noisy and informal. Figure 1 model applies the preprocessing, tokenizing and embedding layers to perform sentiment inference. One is able to write real time and are provided with sentiment prediction and visual explanations with the assistance of explainable AI based on LIME. It is also tested with the assistance of IMDb and twitter data sets which leads to the conclusion that the framework itself is the best and most interpretable simultaneously hence is a comparatively effective and easy to implement in the framework of real-time sentiment analysis in such applications as social media monitoring and customer feedback analysis.

IV. METHODOLOGY

A. Dataset

This paper uses two popular sentiment analysis datasets to measure the performance of the models on various text types and classification problems. The IMDb dataset is

50,000 positive or negative movie reviews, which can be used as a reference point when the author performs binary sentiment classification on longer and more structured text. The Twitter data set consists of 14,640 tweets that have been labeled with three sentiment types, including positive, negative, and neutral, providing a difficult dataset of short, informal, and noisy texts of social media. The two datasets were found in the open repositories and, to facilitate the sound validation of the models and practical usage, they are divided into the training and testing sets.




Figure 2: Class distribution on both datasets

As shown in Figure 2, the class distribution across both datasets is relatively balanced, with the IMDb dataset containing equal proportions of positive and negative reviews, while the Twitter dataset shows comparable representation across positive, negative, and neutral categories. This balanced distribution ensures unbiased model training and reliable evaluation metrics.

B. Data Preprocessing and feature selection

The original social media text raw data were systematically and multi-staged pre-processed into a structured and model-ready text to deal with the inherent noisiness and variability of user-generated information. The initial stages of the preprocessing are normalization of text; here, all the inputs are made to capital letters to improve consistency. As a next step, a thorough clean-up step eliminates URLs, user mentions, hashtags, punctuations, special characters and numeric values, which add no meaning to the sentiment, but create noise. Redundant entries in the texts are then removed to avoid the data leakage and model bias. After that, the NLTK library is used to remove stopwords to eliminate high-frequency words that carry low information content and thus enable the model to concentrate on lexically and sentimentally meaningful words. The tokenization of the cleaned text is done with the WordPiece tokenizer of the Hugging Face Transformers library that divides text into subword units that can be used in any pretrained transformer architecture, e.g. RoBERTa or DistilBERT. Lastly, an exploratory data analysis is done to investigate the distribution of classes, the statistics of review lengths, and the trends in term frequencies so that

the dataset is not skewed. This processed corpus is further divided into training (80) and testing (20) subsets with proportional representation of classes which gives a clean standardised analysis-ready dataset to train robust models on and evaluate them.

C. Model Training

All the models are trained by the early stopping to prevent overfitting. The use of Adam optimizer and binary cross-entropy loss in the case of IMDB and categorical cross-entropy loss in the case of Twitter is utilized. In order to detect generalization, validation is done on a held-out subset.

BiLSTM (Bidirectional LSTM):

BiLSTM operates in both directions and reverses the context both in the past and in the future. It has an embedding layer which transforms words into vectors, BiLSTM layer which works in a bidirectional fashion, dropout layer which does so regularly and dense output layer which classifies using either a sigmoid or softmax activation function.

Mathematical Representation:

$$h_t = \text{BiLSTM}(x_t, h_{t-1})$$

$$\hat{y} = \sigma(W \cdot h_T + b)$$

RoBERTa-BiLSTM:

This bi-skeleton design is the integration of power of RoBERTa transformer to produce rich contextual embedding and BiLSTM layer to acquire sequential correspondence in text. RoBERTa enhances the semantic relationship cognition and BiLSTM takes the word order.

Mathematical Representation:

$$E = \text{RoBERTa}(X)$$

$$h = \text{BiLSTM}(E)$$

$$\hat{y} = \text{Softmax}(W \cdot h + b)$$

DistilBERT+GRU:

DistilBERT is an effective embedding model and GRU is capable of training sequential dependencies at a relatively low level of computational resources in contrast to LSTM. The model can be applied particularly in scenarios where low latency and low computational requirements are required as in real-time applications.

Mathematical Representation:

$$E = \text{DistilBERT}(X)$$

$$h_t = \text{GRU}(E_t, h_{t-1})$$

$$\hat{y} = \sigma(W \cdot h_T + b)$$

V. RESULTS AND DISCUSSION

In this section, a comparative analysis will be performed on the models evaluated for sentiment classification. We analyze the performances of BiLSTM, RoBERTa-BiLSTM, and DistilBERT+GRU on the task of classifying text sentiments as positive, negative, or neutral. The key

results discussed will comprise the effectiveness of each model not only in terms of accuracy but also precision, recall, and F1-score, along with computational efficiency considerations for real-time deployment. The discussion will also address the interpretability of model predictions through LIME-based explanations. Certain assessment measures will be stated both in the following sections and in further elaborate findings.




Figure 3: Confusion matrixes of models on imdb dataset



Figure 4: Confusion matrixes of models on twitter dataset

The confusion matrices of models in figure 3 and 4 indicate that the patterns of classification vary. RoBERTa-BiLSTM shows the most balanced prediction and the least misclassification in both directions, which can be indicative of a good contextual knowledge consideration in the case of IMDB. BiLSTM is more prone to false positives meaning that it is more susceptible to negative aspects. Neutral classes are a significant issue in all models of Twitter analysis, and typically, it is easy to confuse the two. RoBERTa-BiLSTM is depicted to be more effective with negative sentiments and less with neutral cases. These tendencies confirm the fact that hybrid architectures better cope with incompatible language and context, but independent recurrent models have flaws in the multi-faceted sentiment classification, particularly in unclean social media environment.

Table 2: classification report of models on imdb dataset

Model	Accuracy (%)	Precision	Recall	F1-Score
Bilstm	84	0.84	0.84	0.84
Roberta-Bilstm	90	0.90	0.90	0.90
Distilbert+Gru	89	0.89	0.89	0.89

Table 3: classification report of modelson twitter dataset

Model	Accuracy (%)	Precision	Recall	F1-Score
Bilstm	77	0.76	0.77	0.76
Roberta-Bilstm	80	0.79	0.80	0.79
Distilbert+Gru	79	0.78	0.79	0.78

According to Table 2 and Table 3 classification report, in both data sets, RoBERTa-BiLSTM has been the best performer with the highest F1-scores (0.90 IMDb, 0.79 Twitter). The discrepancy in the accuracy and the recall of the binary classification is not marked in all the models except the performance of the neutral classes which is significantly lower when it comes to analyzing Twitter with F1-scores declining to 0.5660. This implies that there is the hard problem of separating the informal text and polarized feelings that is present in the informal text. Computational efficiency Fallaciously, distilled transformers are said to be practical with high performance and efficiency of DistilBERT+GRU. The findings verify the hypothesis that context-rich sentiment analysis tasks where hybrid transformer-recurrent models are very sensitive to semantics and sequence sensitivity, and where standalone models are failing to work well.




Figure 5: Lime explanation for positive prediction




Figure 6: Lime explanation for negative prediction




Figure7: Model accuracy comparison chart

The quantitative comparative analysis in Figure 7 demonstrates a clear advantage of the proposed hybrid architectures over existing sentiment analysis methods. The RoBERTa-BiLSTM model achieves 90% accuracy on the IMDb dataset, outperforming standalone RoBERTa models which typically achieve 87–88% on similar binary classification tasks. This improvement becomes more pronounced on the more challenging Twitter dataset, where the model reaches 80% accuracy compared to the 76–78% range reported for conventional transformer-based approaches. The performance gap between these two datasets aligns with established expectations in the field, where informal social media text presents greater classification difficulty due to noise, ambiguity, and the inclusion of neutral sentiment classes.

The technical superiority of RoBERTa-BiLSTM stems from its dual-mechanism design. The RoBERTa component provides deep contextual understanding through its transformer architecture, excelling at capturing nuanced sentiment indicators, sarcasm, and implied meaning that are prevalent in social media communication. Simultaneously, the BiLSTM layer models sequential dependencies and local syntactic patterns, effectively processing word-order relationships and phrase-level constructions that are essential for accurate sentiment interpretation. This combination allows the model to balance global semantic awareness with local contextual precision, a capability that standalone architectures struggle to achieve.

DistilBERT+GRU offers a compelling alternative for practical applications where computational efficiency is paramount. Through knowledge distillation, this model retains approximately 97% of BERT's language understanding capability while reducing parameter count by 40% and inference time by 60%. Its competitive performance—89% on IMDb and 79% on Twitter—demonstrates that substantial efficiency gains can be achieved without significant accuracy sacrifice, particularly valuable for real-time sentiment monitoring applications and deployment in resource-constrained environments.

Detailed experimental validation confirms the robustness of these findings. Confusion matrices in Figures 3 and 4 reveal that hybrid architectures exhibit more balanced error

distributions than standalone models, with fewer systematic biases in misclassification patterns. LIME-based explanations in Figures 5 and 6 provide transparent insights into model decision-making, showing that hybrid models consider more contextual features per prediction and demonstrate greater sensitivity to linguistic subtleties like negation and intensification. Additionally, the accuracy comparisons in Figure 7 further validate the performance advantages of hybrid models over existing approaches. These empirical observations from Figures 5, 6, and 7 directly explain the performance advantages observed in quantitative metrics.

The practical implementation of these models in a functional web application provides additional validation through real-world testing. User interactions with the system have confirmed laboratory results, with only minimal deviation in performance metrics when processing diverse, unseen text samples. This deployment not only verifies the experimental findings but also demonstrates the system's readiness for practical applications in social media monitoring, customer feedback analysis, and other sentiment tracking scenarios where accuracy, interpretability, and usability are equally important considerations.

VI. CONCLUSION AND FUTURE WORK

This research successfully demonstrates the efficacy of hybrid deep learning architectures for advanced sentiment analysis. The integration of transformer-based contextual modeling with recurrent sequential processing yields a system capable of nuanced interpretation across both formal reviews and informal social media text. The incorporation of explainable AI techniques provides essential transparency, bridging the gap between complex model internals and user comprehension. The developed web application operationalizes this research, offering an accessible and trustworthy tool for real-time sentiment evaluation. Future work will focus on expanding the system's linguistic and functional scope, including integrating multilingual support to process global social media content and refining the model to handle more granular sentiment spectrums such as mixed or ambivalent emotions. Enhancements in computational efficiency will be pursued to enable large-scale, real-time analytics, while advanced transfer learning and domain adaptation techniques will improve robustness across different industries and contexts. Furthermore, the explainability framework will be extended with complementary methods to provide even deeper insights into model reasoning, ensuring the system remains both state-of-the-art and user-centric in evolving applications.

REFERENCES

- [1] S. Kiritchenko and X. Zhu, "Sentiment analysis meets explainable artificial intelligence: A survey on explainable sentiment analysis," *IEEE Trans. Affect. Comput.*, vol. 15, no. 3, pp. 837-846, Jul.-Sep. 2024, doi: 10.1109/TAFFC.2023.3296373.
- [2] Q. Lu, X. Sun, Y. Long, Z. Gao, J. Feng, and T. Sun, "Sentiment analysis: Comprehensive reviews, recent advances, and open challenges," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 35, no. 11, pp. 15092-15112, Nov. 2024, doi: 10.1109/TNNLS.2023.3294810.
- [3] H. Huang, A. A. Zavareh, and M. B. Mustafa, "Sentiment analysis in e-commerce platforms: A review of current techniques and future directions," *IEEE Access*, vol. 11, pp. 90367-90382, Aug. 2023, doi: 10.1109/ACCESS.2023.3307308.
- [4] Y. Aliyu, A. Sarlan, K. U. Danyaro, A. S. B. A. Rahman, and M. Abdullahi, "Sentiment analysis in low-resource settings: A comprehensive review of approaches, languages, and data sources," *IEEE Access*, vol. 12, pp. 66883-66909, May 2024, doi: 10.1109/ACCESS.2024.3398635.
- [5] A. A. Alharbi, M. A. Alshayea, and M. A. Alshayea, "Aspect-based sentiment analysis: A comprehensive review and open research challenges," *IEEE Access*, vol. 13, pp. 65138-65182, Mar. 2025, doi: 10.1109/ACCESS.2025.3555744.
- [6] M. A. AlZubi and M. Al-Ayyoub, "Challenges and issues in sentiment analysis: A comprehensive survey," *IEEE Access*, vol. 11, pp. 69626-69642, 2023, doi: 10.1109/ACCESS.2023.3293041.
- [7] V. S. Anoop, T. K. Ajay Krishnan, A. Daud, A. Banjar, and A. Bukhari, "Climate change sentiment analysis using domain specific bidirectional encoder representations from transformers," *IEEE Access*, vol. 12, pp. 114912-114922, Aug. 2024, doi: 10.1109/ACCESS.2024.3441310.
- [8] C. Liu, A. Arulappan, R. Naha, A. Mahanti, J. Kamruzzaman, and I.-H. Ra, "Large language models and sentiment analysis in financial markets: A review, datasets, and case study," *IEEE Access*, vol. 12, pp. 134041-134061, Aug. 2024, doi: 10.1109/ACCESS.2024.3445413.
- [9] A. Albladi, M. Islam, and C. Seals, "Sentiment analysis of Twitter data using NLP models: A comprehensive review," *IEEE Access*, vol. 13, pp. 30444-30468, Feb. 2025, doi: 10.1109/ACCESS.2025.3541494.
- [10] J. Li et al., "Multimodal sentiment analysis—A comprehensive survey from a fusion methods perspective," *IEEE Access*, vol. 13, pp. 64556-64583, Mar. 2025, doi: 10.1109/ACCESS.2025.3554665.
- [11] M. M. Hossain, S. Sanjara, M. S. Hossain, S. Chaki, M. S. Rahman, and A. B. M. S. Ali, "Revolutionizing mental health sentiment analysis with BERT-Fuse: A hybrid deep learning model," *IEEE Access*, vol. 13, pp. 85428-85446, May 2025, doi: 10.1109/ACCESS.2025.3568340.
- [12] N. Liu and J. Zhao, "Recommendation system based on deep sentiment analysis and matrix factorization," *IEEE Access*, vol. 11, pp. 16994-17001, Feb. 2023, doi: 10.1109/ACCESS.2023.3246060.